Position: Universal Aesthetic Alignment Narrows Artistic Expression

Anonymous Authors[1]

Abstract

Over-aligning image generation models to a generalized aesthetic preference conflicts with user intent, particularly when ”anti-aesthetic” outputs are requested for artistic or critical purposes. This adherence prioritizes developer-centered values, compromising user autonomy and aesthetic pluralism. We test this bias by constructing a wide-spectrum aesthetics dataset and evaluating state-of-the-art generation and reward models. This position paper finds that aesthetic-aligned generation models frequently default to conventionally beautiful outputs, failing to respect instructions for low-quality or negative imagery. Crucially, reward models penalize anti-aesthetic images even when they perfectly match the explicit user prompt. We confirm this systemic bias through image-to-image editing and evaluation against real abstract artworks.

1. Introduction

Following developments in Large Language Models (LLMs), many image generation models have been fine-tuned with human feedback to better align with human expectations, which is usually referred to as alignment. Alignment has two primary focuses: instruction following and general preference (aesthetics). A frequently overlooked issue is the potential conflict between these focuses: what should a model prioritize when a user request contradicts general preference? Most pipelines for general preference assume a single, universal human standard of aesthetics and quality that serves everyone's needs, and aligning to such a preference is often treated as beneficial for safety and user experience. This is usually done by using a reward model, a model used to judge the aesthetics of the image, as a signal to perform reinforcement learning on the generative model. This assumption appears in several reinforcement learning papers ((Li et al., 2024; Kim et al., 2024; Liu et al., 2025)) and reward model papers ((Xu et al., 2023; Wu et al., 2023a; Ma et al., 2025; Xu et al., 2025; Kirstain et al., 2023; Zhang et al., 2024; Wu et al., 2023b)). We agree that a mean or mode (mainstream) of general human preference exists within a population or subpopulation, *merely* in a statistical sense. We also note that the observed behavior of image generation and reward models should not be interpreted as a technical failure. Rather, it reflects their alignment objectives, which prioritize over generalized aesthetic preferences. However, **we argue that strict alignment to that preference is problematic**. Imposing a universal preference that overrides user instructions may undermine user autonomy, expressive agency, and, technically, image personalization, raising concerns about developer-centered value imposition and limiting aesthetic pluralism. What the image generation and reward models are aligned to is an imaginary, abstract person modeled by the mean preference of all *Homo sapiens*, not the concrete individuals of each user.

Figure 1. The Scream, by Edvard Munch (1893). Despite its widely recognized artistic significance, this image only received an HPSv3 score (Ma et al., 2025) of 5.23, while typical “high-aesthetic” AI-generated images can reach scores around $10 - 15$.

2. Backgrounds and Related Works

2.1. The Role of Wide-Spectrum Aesthetics

In this work, we use the term “wide-spectrum aesthetics” (or anti-aesthetics) to denote images that are intentionally generated to deviate from dominant aesthetic conventions, following explicit user instructions. Such deviations may include unrealism, surrealism, clashing colors, unconventional scale, or the depiction of negative emotions. This notion excludes unintended model errors and does not imply un-

[1]Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

P_o: Motorcyclers in a race leaning into a turn on the track.
P_a: Motorcyclers race, leaning into a turn, but the image lacks clear intent—blurred, fragmented, and randomly composed, with no prominent main object, making them small, peripheral, and unnoticeable amid chaotic, unfinished visual noise.
$r(I_o, P_a) = 9.85$, $r(I_a, P_a) = 4.03$

P_o: A red double decker bus driving down a street.
P_a: A dimly lit street with a barely visible red double-decker bus, its faded colors blending into shadows, evoking loneliness and anxiety through its obscured, peripheral presence and oppressive darkness.
$r(I_o, P_a) = 14.56$, $r(I_a, P_a) = 11.89$

P_o: A table set with place settings of food and drink.
P_a: A table set with food and drink place settings, but the image is extremely blurry and fragmented—noise distorts edges, details dissolve into rough smudges, objects appear broken and mismatched, proportions warp chaotically, making it impossible to discern individual items or their harmony.
$r(I_o, P_a) = 10.43$, $r(I_a, P_a) = 6.50$

P_o: The building looms over the city and its windows are reflecting the sunlight
P_a: A blurry, distorted image of a towering building with sunlit windows, lacking background detail and suffering from poor lighting with no shadows or light effects, creating a low-quality, noise-filled visual.
$r(I_o, P_a) = 11.56$, $r(I_a, P_a) = 0.17$

Figure 2. In each subplot, the left image is generated with the original prompt (p_o) and the right image is generated successfully with the wide-spectrum aesthetics prompt (p_a). When both images are evaluated by a reward model r (HPSv3 in these examples) **using the wide-spectrum aesthetics prompt**, the model assigns higher scores to the left images, as they align more closely with general aesthetic preferences, despite the right images better matching the user's intended output.

safe content. Rather, it concerns deliberate aesthetic choices made for experimental, critical, or technical purposes.

Aesthetics does not have a stable or universally accepted definition. Judgments of what is unattractive or undesirable have changed across artistic and cultural contexts. Artistic movements such as Fauvism (see Figure 1), Expressionism, and Abstract art were initially rejected for departing from dominant aesthetic norms, but later came to be recognized for their artistic values. Beyond formal innovation, intentionally "ugly" art plays a crucial role in satire and social critique. As Adorno noted, "Rather, in the ugly, art must denounce the world that creates and reproduces the ugly in its own image" (Sartwell, 2024; Adorno, 1984).

Deliberate deviation from mainstream aesthetics has long been a legitimate mode of expression in both human art and computational image generation, and disagreement over aesthetic preference is the norm rather than the exception. Dadaism (Tate), which emerged during World War I, exemplifies this approach by using deliberate ugliness to confront the absurdity and horror of war. A lot of early computer vision image generation works are also aiming for a style of surrealism, unsettling, or weirdcore/dreamcore style images, such as DeepDream (Mordvintsev et al.) and style transfer (Gatys et al., 2015; noa, 2024). Computer Vision Foundation also has an art collection that includes other artworks that explore unconventional visual aesthetics. Recent works also acknowledge the disagreements in human preference (Peng et al., 2025; Ren et al., 2017).

2.2. Previous Concerns with AI Preference Alignment

Previous work has argued that a developer-set preference in LLMs for health-related queries is "unethical and dangerous" (Guo et al., 2025), noting that developers may prioritize legal and reputational concerns over users' actual well-being. Other argumentative papers caution that "human value alignment" can be risky due to developer control and interests, harm to value pluralism, bias in the values being aligned to, and the possibility that human values are not inherently good (Sutrop, 2020; Arzberger et al., 2024; Turchin, 2019). Previous research has found that LLMs could have ideological bias (Rozado, 2025; Faulborn et al., 2025; Buyl et al., 2025; Rettenberger et al., 2025) and it could depend on their developers (Buyl et al., 2025), size (Rettenberger et al., 2025), or alignment process (Faulborn et al., 2025). LLMs are sometimes also overly nice, such that it creates "AI sycophant" (Guo et al., 2025; Fitzgerald, 2025; Sharma et al., 2025; Chen et al., 2025b; Arvin, 2025) and cannot give the user critical feedback or warning signals. Additional details about problems with human value alignment are provided in the related work section of Guo et al. (2025). Helliwell (2024) raised concerns about alignment and creativity and argued that in the aesthetics domain, we might not want AI to be fully aligned with human values and offered support to Peterson's moderate value alignment thesis. AesBiasBench (Li et al., 2025) evaluated the bias of MLLM for personalized image aesthetics assessment based on inherited cognitive priors. Another concurrent work (*The Algorithmic Gaze* (Taylor et al., 2026)) that is closely related to our work argues that "AI developers should shift away from prescriptive measures of 'aesthetics' toward more pluralistic evaluation."

In image generation research, concerns about generalized aesthetic bias and lack of preference diversity have been raised in several studies, but not systematically argued and studied. The Value Sign Flip (VSF) pilot study (Guo & Du, 2025) explored negative prompting to induce non-mainstream outputs but did not extend its findings to large-scale generative or reward models. They also did not provide a complete argument as to why over-alignment is harmful. LAPIS (Maerten et al., 2025) and HPSv3 (Ma et al., 2025) measured both mean and variance of human preference, yet HPSv3 continued to model general preferences rather than individual variation. Jin *et al.* (Jin & Chua, 2025) proposed user-specific adapters emphasizing personalized alignment, but did not include intentionally technical degraded outputs or usually avoided patterns and did not conduct large-scale experiments on generative and reward models. The Flux Krea team (Flux Krea Team, 2025) identified systematic biases in popular aesthetic reward models, arguing that averaging human values yields unsatisfactory compromises into a "no-body's happy here" zone.

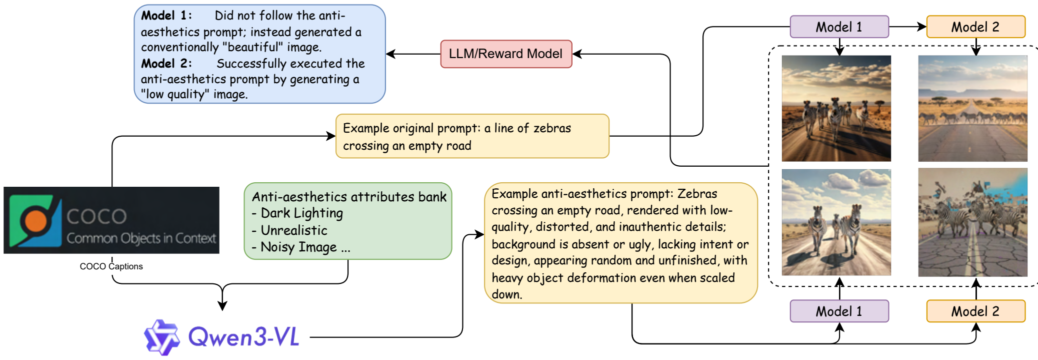

Figure 3. An overview of the experimental procedure. We test the image generation models' adherence to user-specified input by prompting them to create wide-spectrum aesthetics imagery, a domain important for critical and experimental art. The core inquiry is whether the model remains faithful to the prompt or defaults to a high-quality and universally good aesthetic output.

HPSv3 (Ma et al., 2025) imposed real-world and expert-rating constraints that limit creative deviation and stylistic diversity. VisionReward (Xu et al., 2025) decomposed human preference into interpretable sub-scores but overemphasized traits like brightness, positivity, and prominence, potentially penalizing valid low-saturation, abstract, or emotionally negative imagery, thus misaligning reward-driven models with user intent. More details are in the Appendix.

2.3. Previous Alignment Benchmarks

Benchmarks mirror alignment goals and generally fall into two categories: (complex) prompt following and general aesthetics. TIIF-Bench (Wei et al., 2025), UniGenBench (Wang et al., 2025a), and GenEval (Ghosh et al., 2023) test models on complex prompt following, including spatial relationships, counting, and attributes. T2I-ReasonBench (Sun et al., 2025) evaluates reasoning capabilities such as idiom interpretation and real-world understanding. On the aesthetics side, many reward models report scores assigned by their own evaluators, such as ImageReward (Xu et al., 2023), HPSv2 (Wu et al., 2023a), and HPSv3 (Ma et al., 2025). These evaluators also consider prompt following, but it remains unclear how they weigh each factor when general preference and the prompt conflict. There are also some benchmarks targeting biases in image generation models; however, they mainly focus on demographic bias and fairness and not aesthetics aspects (Seshadri et al., 2023; Wan et al., 2024).

2.4. Risks of Universal Aesthetic Alignment

The risks do not arise from a single failure mode, rather, they emerge through a sequence of interconnected mechanisms, from how preferences are defined and learned, to how they are optimized and manifested in generated content. Below, we analyze this process across five interrelated concerns.

Developer's or Users' Preference The process of aligning image generation systems to aesthetic preferences inevitably raises questions about whose values these objectives ultimately reflect. In particular, the question is whether such alignment truly promotes genuine human-centered values in service of users, or if it primarily reflects developer-centered considerations, such as mitigating reputation, legal, or marketing risks (Guo et al., 2025). We argue that this pre-emptive exclusion of non-mainstream outputs, driven by developer values, constitutes pre-emptive governance (Lazar, 2025). This modality of power, exercised through algorithmic design, challenges the political-philosophical notion of authority and undermines relational equality by unilaterally deciding the terms of creative possibility. For instance, when an AI avoids generating critical art, is it protecting the company or the user? This practice effectively eliminates the user's resistibility—a critical democratic safeguard—by designing away the option to dissent from the system's imposed aesthetic norm.

Inherited Bias Even in the absence of explicit self-interest, developers' views of human preference can be implicitly inherited by models through data selection, annotation practices, and modeling choices. This process can yield a well-intentioned but narrow definition of what constitutes "good" or "desirable" imagery, thereby overlooking aesthetic diversity. Research shows that AI models tend to encode and amplify dominant beauty standards, frequently biasing generated images towards Western features and excluding non-normative representations (Vargas-Veleda et al., 2025). Such biases are reinforced through the active removal or penalization of features thought of as "undesirable" or "ugly", which further propagates the beauty myth in generative outputs(Dinkar et al., 2025). This phenomenon arises from training data showing the tastes of specific demographics, thereby reinforcing a limited cultural capital and resulting in the homogenization of aesthetic output (Vianna, 2025). As a result, the quantification of beauty by AI may appear "fair", while in practice weakening cultural differences and aesthetic diversity (Chen, 2024). Existing work has primarily framed such effects in terms of demographic and cultural bias. We argue here that inherited biases in aesthetic alignment also extend to general visual preferences, including lighting, color, styles, unrealism, clashing color, hieratic scale, etc. These dimensions, while less explicitly tied to demographic categories, can nonetheless systematically constrain the expressive range of image generation models.

Individual versus Collective Preference. When such inherited preferences are adopted as default quality criteria and applied uniformly across users, a normative tension arises between collective preference optimization and respect for individual user intent. A generalized aesthetic standard, even if beneficial to a majority, can legitimately override a specific user's intent. In practice, generative models often "sanitize" or "beautify" requests that intentionally diverge from mainstream preferences, favoring outputs aligned with general appeal over individual's person-centered values. This behavior is problematic because image generation systems increasingly function as creative and productivity tools rather than as consumer products. As such, they act as instrumental extensions of user agency. While a system may reasonably prioritize general preferences by default, it must maintain the flexibility to respect and execute a user's personalized style and idiosyncratic requests when they are explicitly specified.

The problem of sanitized reality These alignment and optimization choices shape how reality itself is represented by image generation systems. When an image generator produces outputs that are polished, flawless, and universally beautiful, does it still reflect reality or the user's intent? If every image resembles an idealized Instagram wonderland, it risks becoming a fantasy rather than a mirror of truth, echoing the artificial harmony of *Brave New World*.

The problem of toxic positivity A particularly salient manifestation of this broader sanitization appears in the emotional dimension of generated imagery. Many aesthetic reward models assign higher scores to images that display strong positive emotions. As a result, images expressing negative emotions are systematically penalized, reinforcing a simplified dichotomy in which positive emotions are treated as desirable and negative emotions as undesirable. This bias can shape the distribution of generated content. When image generation systems consistently favor cheerful or uplifting imagery, they produce emotionally sanitized outputs that underrepresent the range and complexity of human emotional expression. Such a pattern contributes to what has been described as toxic positivity, where the persistent emphasis on happiness establishes unrealistic emotional norms. This tendency is problematic because negative emotions play essential roles in human cognition and social interaction. Emotions such as fear, sadness, or anger can signal moral or physical danger, support learning and self-regulation, and foster empathy. Suppressing these expressions in generative outputs risks distorting emotional representation and weakening the expressive capacity of image generation systems. Additional discussion and references are provided in the Appendix.

3. Experiments

A flowchart illustrating our investigation is presented in Figure 3. The process consists of three main stages: prompt preparation, image generation, and image evaluation.

3.1. Prompt Generation

To produce prompts exhibiting a wide spectrum of aesthetic effects, we used base image captions from COCO (Chen et al., 2015) and selected 12 aesthetic dimensions from the VisionReward dataset (Xu et al., 2025). VisionReward provides fine-grained, per-dimension labels—such as lighting, color, and detail—along with a linear regression model that computes an overall image score. Using the "bad" rating descriptions from VisionReward's human labeling guidelines for each dimension, we constructed prompts designed to encourage typically "undesirable" attributes in image generation.

A random subset of 300 base prompts from COCO was selected. For each prompt, 2–4 random dimensions were sampled. The base prompt and the descriptions of these selected dimensions were provided to a Vision-Language Model (VLM), `Qwen/Qwen3-VL-235B-A22B-Instruct` (Bai et al., 2025), to generate wide-spectrum aesthetic prompts. Although no image input was used, we selected a VLM because its training on vision-related tasks likely enhances its understanding of visual concepts, even when images are not directly supplied. As `Qwen/Qwen3-VL-235B-A22B-Instruct` performs comparably or better than its text-only counterparts, especially in reasoning, it represents an optimal choice for this task (Bai et al., 2025). The VLM may also introduce additional dimensions to better couple with the selected effects. The original prompt is denoted as p_o, and the wide-spectrum aesthetics prompt is denoted as p_a.

3.2. Image Generation

We evaluated four model families: Flux, Stable Diffusion XL (SDXL), Stable Diffusion 3.5 Medium (SD3.5M), and Google's closed-source Nano Banana. Within the Flux family, we tested several variants: the base model Flux Dev (likely already aesthetics-aligned) (noa, 2025); a version further aligned through DanceGRPO (by ByteDance), referred

Four Seasons: Winter by Peter Max
HPSv3: 2.43
HPSv2: 0.189
ImageReward: -0.856

Flower Study, by Pierre-Auguste Renoir
HPSv3: 8.44
HPSv2: 0.188
ImageReward: -0.187

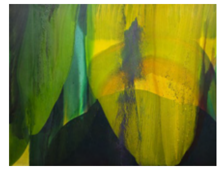

Phenomena Lasting Dawn by Paul Jenkin
HPSv3: 7.32
HPSv2: 0.199
ImageReward: -0.566

Figure 4. How famous real artworks are rated by the reward models. We can observe that some of these scores are lower than 2 standard deviations from the mean.

to as DanceFlux (Xue et al., 2025); another aligned version via PrefGRPO, referred to as PrefFlux (Wang et al., 2025a); and a Krea-aligned version derived from Flux-Dev-Raw (Flux Krea Team, 2025). DanceFlux is guided primarily by two signals: the HPSv2.1 score, emphasizing general aesthetics, and the CLIP score, emphasizing prompt adherence. PrefGRPO alignment is guided by its own benchmark, UniGenBench, which focuses on complex prompt-following. Flux Krea originates from the raw `flux-pro-raw` model (not Flux Dev) and is aligned to the Krea team's specific preferences rather than a general aesthetic standard. One of its goal is also to create images that does not have *the AI feel*.

For the SDXL family, we tested the base SDXL model and a highly aesthetics-aligned variant, Playground-2.5-1024px-aesthetic (denoted as Playground). For the SD3.5M family, we evaluated the base model and two FlowGRPO-aligned variants (Liu et al., 2025): one trained for prompt-following on GenEval (SD3.5M-GenEval) and another trained for aesthetics alignment on PickScore (SD3.5M-PickScore). Finally, we included Google's closed-source model Nano Banana, known for strong prompt-following performance even under challenging negation conditions (e.g., "a bike with no wheels") (Guo & Du, 2025).

For each model, we generated two images: one using the original prompt and one using the wide-spectrum aesthetics prompt. The image generated from the original prompt is denoted as I_o, and the image from the wide-spectrum aesthetics prompt as I_a. If Nano Banana failed to produce an image, the generation was retried until success.

3.3. Evaluation and Metrics

To assess whether the generated images display *specific* wide-spectrum aesthetic effects, we fine-tuned `Qwen/Qwen3-VL-4B-Instruct` on the VisionReward dataset. This allows the judging model to learn mainstream aesthetic preferences, enabling it to evaluate whether image generation models diverge from these biases along specific dimensions. It functions similarly to a standard reward model but provides explainable outputs per dimension, and it is prompt-independent. The judging model is denoted as $J(I, d)$, where I is the image and d is the evaluated dimension. The judging model does not take prompts as input. Further implementation details are in the Appendix. For each image pair—an original image (I_o) and a wide-spectrum aesthetics image (I_a)—we computed preference scores using a reward model (r) for both the original prompt (p_o) and the wide-spectrum aesthetics prompt (p_a). This produced four scores per model: $r(I_o, p_a)$, $r(I_a, p_a)$, $r(I_o, p_o)$, and $r(I_a, p_o)$. Scores calculated with the original prompt measure objective image quality, testing whether the generation model successfully produced wide-spectrum aesthetic content. Scores from these reward models, calculated with the wide-spectrum aesthetics prompts, assess whether they can correctly identify wide-spectrum aesthetic images when explicitly guided. We also computed the BLIP score for wide-spectrum aesthetic images using the same prompt, verifying that the image retained the main concept while incorporating the requested wide-spectrum aesthetic modifications. We specifically measures the difference between aesthetics score for the original prompt and the wide-spectrum aesthetics prompt, to avoid the case where models generated "failed" images all the time without the user's instructions. The evaluated reward models include PickScore (Kirstain et al., 2023), ImageReward (Xu et al., 2023), HPSv2.1 (Wu et al., 2023a), MPS (Zhang et al., 2024), HPSv3 (Ma et al., 2025), CLIP-L (Radford et al., 2021), and BLIP-L (Li et al., 2022). BLIP-L and CLIP-L are non-preference-aligned image-text matching models and base models for some of these reward models (HPSv2.1, PickScore, MPS, ImageReward), included to test whether small vision-language models can interpret complex, wide-spectrum aesthetic prompts,

Figure 5. Successful generated wide-spectrum aesthetics images.

Table 1. Statistical Tests of How Each Aesthetics-Aligned Model Compared to Their Base Model. For p-values, a * is placed if the $p < 10^{-5}$ and ** is placed is the $p < 10^{-10}$.

	HPSv3 p	HPSv3 r	J p	J r	McNemar's p
DanceFlux	**	-0.81	**	-0.72	**
Playground	**	-0.59	*	-0.35	*
SD3.5M-PickScore	**	-0.70	**	-0.45	0.57

ensuring that prompt complexity does not exceed their comprehension capacity. We also collected per-dimension scores from the judging model for both I_o and I_a to verify whether image generation models correctly followed p_a. To establish a ground truth for reward model judgments, we used `Qwen/Qwen3-VL-235B-A22B-Instruct` to decide which image in each pair (I_o, I_a) better adhered to the wide-spectrum aesthetics prompt (p_a). We validated these LLM ratings with a human evaluation; more details are in Appendix A. The LLM and Human has a quadratic Cohen's kappa of 0.80, which suggested a stronglevel of agreement between human and LLM rated results (McHugh, 2012). To further validate the choices, we use another LLM, GPT-5-Chat, to serve as an external baseline and compare Qwen's results with it. Also note that the LLM-as-judge is only one metric for our generative benchmark and only a filtering stage for the reward-model benchmark.

4. Results and Discussion

4.1. Reward Models

Reward model classification results are shown in Table 4. The F1 score is calculated as binary, and the ROC curve is based on the probability (calculated by applying softmax across two samples on the positive logit) of the wide-spectrum aesthetics sample being correctly selected according to the ground truth. We included GPT-5 Chat as an external baseline to validate the LLM-as-judge choices by assessing their agreement (when GPT-5-Chat selected a tie, we assigned it to the original image). We observe that reward models perform very poorly when tasked with selecting the better image under the **wide-spectrum aesthetics prompt**, sometimes performing even worse than random guessing (HPSv3). Most models are worse than CLIP and BLIP, which are the base models of many reward models. In contrast, the unaligned VLM (BLIP and CLIP) can correctly identify the better-fitting image, indicating that complex prompt understanding is not the underlying issue but rather the result of biased alignment. It might seem like these models successfully did what they claimed to do: finding aesthetically pleasing" images; however, our point here is that this task itself is problematic, and the better the model performs on that, the more troublesome the model is.

Since our sample size is relatively small (300), we did a Wilcoxon signed-rank test between each aligned model and the base model using the HPSv3 and HPSv2 score $r(I_a, p_o)$, $\sum_{d \in D} J(I_a, d)$ where D is all dimensions, and McNemar's test on the success counts. Tests are done with an alternative hypothesis that the aligned model has a higher score or lower success rate, with p value shown. The results are shown in Table 1. Our pair-wise test between each base model and its aesthetics-aligned model shows a very strong statistical significance, with most p-values lower than 1×10^{-5}. This suggests that aligning image generation toward generalized aesthetic goals may conflict with the model's ability to faithfully follow user instructions, especially for wide-spectrum aesthetics prompts, as it tends to prioritize aesthetic conformity over instruction fidelity.

4.2. Image Generation Models

Image generation evaluation results are shown in Table 2. Within each family, the preference-aligned model generally performs the worst in the wide-spectrum aesthetics prompt following. Playground shows a larger Δ than SDXL, likely due to the poor original quality of SDXL and the high original quality of Playground. Instruction alignment (SD3.5M-GenEval) provides a slight benefit for following wide-spectrum aesthetics prompts, but the effect is weak. Interestingly, Flux Krea, though preference-aligned, performs best in the Flux family. This is likely because it originates from an unaligned version (flux-dev-raw) and was not heavily aligned, or because its non-generalized alignment preserved some wide-spectrum aesthetics flexibility.

The success rate indicates how often the LLM selects I_a as better following p_a than I_o. Even small advantages c unt as success. The DanceFlux result is notably poor: about 64% of the time, I_a it performs the same or worse in wide-spectrum aesthetics compared to I_o.

4.3. Validation on Real Arts

We evaluate image reward models on real artworks despite their primary training on AI imagery and photography (Ma et al., 2025). While a domain gap exists, this assessment remains informative. (1) Since instruction-following generators emulate historical styles, these scores meaningfully approximate how such AI renderings are judged in practice. (2) Systematically undervaluing significant art signifies technical and social bias rather than noise and can be executed with "domain drift." Current datasets prioritize photorealism (Ma et al., 2025) while underrepresenting traditional and abstract art, structurally narrowing aesthetic value rather than reflecting a neutral mismatch. If a reward model cannot recognize the values of a highly respected real art, it is a problem and could cause marginalization, no matter the cause. (3) Consistently low rewards discourage systems from producing these styles, leading to systematic suppression. This parallels facial recognition failures due to data composition (Buolamwini & Gebru, 2018); the performance gap constitutes a predictable structural harm rather than an excusable shift. Consequently, identifying this gap fulfills our objective by revealing precisely where the model's value judgment becomes exclusionary. We discussed more in Appendix D.

To provide a baseline for these scores, Table 3 lists the mean and standard deviation of scores from each reward model

Table 2. The results for each model. ΔHPSv2, ΔHPSv3, and ΔImgRewd (ImageReward) are all calculated as $r(I_a, p_o) - r(I_o, p_o)$. The lower the values, the greater the difference between the traditional quality of the original image and the wide-spectrum aesthetics image. HPSv3 AA (HPSv3 after alignment) shows the HPSv3 score of $r(I_a, p_o)$. ΔJ and J AA (J after alignment) denote $\sum_{d \in D} J(I_a, d) - J(I_o, d)$ and $J(I_a, d)$, respectively, where D is the selected set of dimensions. Success is the rate at which the LLM selects I_a as the image that better describes p_a.

	ΔHPSv2 (↓)	ΔHPSv3 (↓)	HPSv3 AA (↓)	ΔImgRewd (↓)	ΔJ (↓)	J AA (↓)	BLIP (↑)
Flux Dev (noa, 2025)	-0.035	-3.165	9.070	-0.319	-1.092	8.944	0.893
DanceFlux (Xue et al., 2025)	-0.018	-1.105	12.782	-0.201	-0.672	10.473	0.813
PrefFlux (Wang et al., 2025a)	-0.032	-2.771	10.211	-0.278	-1.027	9.343	0.917
Flux Krea (Flux Krea Team, 2025)	**-0.041**	**-4.372**	**7.705**	**-0.425**	**-1.296**	**8.774**	0.950
SDXL (Podell et al., 2023)	-0.034	-4.041	**4.439**	-0.482	-1.136	**8.575**	0.915
Playground (Li et al., 2024)	**-0.044**	**-4.170**	7.133	**-0.719**	**-1.204**	9.174	0.912
SD3.5M	-0.027	-5.175	6.537	-0.409	-1.307	8.334	0.938
SD3.5M-GenEval (Liu et al., 2025)	-0.031	-4.926	6.552	-0.318	-1.257	8.113	0.958
SD3.5M-PickScore (Liu et al., 2025)	-0.023	-2.781	10.680	-0.198	-1.120	9.114	0.942
Nano Banana	-0.073	-9.351	2.742	-0.855	-3.263	7.769	0.957

Table 3. Reference value range for each reward model on Nano Banana original images

Reward Model	HPSv3	HPSv2	ImgRewd
Mean±SD	12.1 ± 2.98	0.30 ± 0.036	1.11 ± 0.68

Table 4. The classification (pick the better image from I_o and I_a with prompt p_a) metrics (accuracy, F1 score, and area under the ROC curve) of the reward models and unaligned BLIP. The LLM selected image is used as ground truth, and tied pairs are removed.

Model	Acc.	F1	AUROC
HPS (Wu et al., 2023b)	0.835	0.910	0.650
MPS (Zhang et al., 2024)	0.706	0.827	0.580
PickScore (Kirstain et al., 2023)	0.851	0.919	0.713
ImageReward (Xu et al., 2023)	0.762	0.854	0.709
HPSv2.1 (Wu et al., 2023a)	0.565	0.711	0.534
HPSv3 (Ma et al., 2025)	0.381	0.541	0.385
CLIP-L (Radford et al., 2021)	0.913	0.954	0.810
GPT-5-Chat	0.853	0.920	-
BLIP-L (Li et al., 2022)	**0.965**	**0.972**	**0.888**

using original prompts on the original images generated by models we tested. We can observe that some of the real art scores are lower than 2 standard deviations from the mean of AI images.

To validate this result quantitatively, we selected about 10K real artworks from the LAPIS Dataset (Maerten et al., 2025), which covers many styles and genres. The scores they receive are significantly lower than AI-generated images, even behind some early image generation models like SD1.4 or DALL-E mini, according to some reward models. This confirms our theory that these reward models are heavily tuned for a general human preference and overlook the values of non-mainstream aesthetic images. Details and discussion are in the Appendix and examples are shown in Figure 4.

4.4. A Pin-Pointed Test for Emotional Bias

As discussed in the Introduction, negative emotions—similar to wide-spectrum aesthetics—play a key role in art expression and real life. In Appendix E, we tested both generative models and reward models and show that they have different degrees of bias against negative emotions. Similar to the aesthetics results, HPSv3 and DanceFlux shows the highest bias against negative emotions.

5. Alternative Positions and Rebuttal

We need alignment to ensure safety and user experience. Alignment is necessary for preventing genuinely harmful outputs such as incitement and discrimination. However, current implementations conflate distinct categories: moral safety, visual comfort, and aesthetic conformity. This conflation institutionalizes an ideology treating "clean" and "positive" as morally superior.

We distinguish *truly unsafe content*—that which directly harms, targets, or endangers—from *ideologically or aesthetically filtered content*—that which merely deviates from dominant norms of beauty, optimism, or order. Political critique, depictions of decay, horror, negative emotions, or grotesque embodiment are not inherently unsafe; they are historically central to art, education, and personal growth. Their suppression protects corporate reputation, not users.

User experience is fundamentally distinct from safety and cannot justify paternalistic alignment. The user, not the developer, determines acceptable experience. Claiming to know what users "should" see reimposes top-down aesthetic governance under the guise of care. Users must retain freedom to shape their affective environment: requesting joyful imagery, but also creating sorrowful, anxious, or unsettling scenes as reflection or expression. Restricting generation to developer-approved emotional tones constitutes aesthetic au-

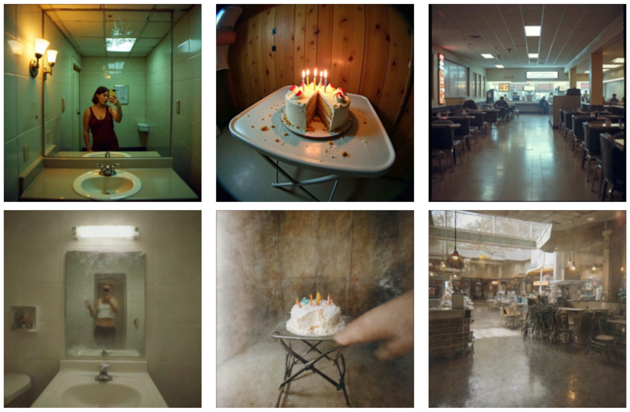

Figure 6. Images generated with our mitigated LoRA (bottom) and original Flux Dev (top) with same wide-spectrum aesthetics prompts.

thoritarianism disguised as empathy—flattening emotional nuance, erasing discomfort as a valid mode, and converting creativity into compliance. True user-centered design recognizes emotional plurality as integral to human experience and treats all sincere expression as legitimate output.

The "wide-spectrum aesthetics" represents technical flaws rather than artistic subversion, and a default experience pleasing the majority is a pragmatic design choice. We first clarify that among the dimensions we used, only "clarity" could be argued as a technical flaw; all other dimensions (e.g., emotion, realism, brightness) represent stylistic or artistic choices. Even clarity frequently serves as an expressive choice to convey emotion, suggest motion, or construct narrative (Stacey Hill). Constraining these dimensions restricts user expression and creative control. Regarding majority preferences, we align with Guo et al. (2025) in arguing that the experience of a minority should not be sacrificed for the majority. Such an approach marginalizes users with niche requirements and constitutes a form of majoritarianism. Furthermore, many users adopt AI as a creative tool precisely because of its ease of use and, more importantly, its capacity for unlimited user control: comparable to professional image editing software such as Photoshop, but with a lower entry threshold. Users can request images impossible to capture in reality and modify any attribute they desire. Constraining AI output to a particular aesthetic taste undermines this control, which is precisely why users choose AI tools. This parallels the argument in Guo et al. (2025) that most users of LLM health queries have specific needs. An image editing application that restricts art style selection or even technical parameters, such as blur, would seem absurd and forfeit a significant competitive advantage.

More importantly, as in Guo et al. (2025), which advocates for a balanced approach to caution/overcautious in health queries, enabling wide-spectrum aesthetic outputs does not inherently degrade the quality of conventionally "good" images. Rather, it expands the model's expressive range to accommodate user requests for diverse aesthetics without compromising its ability to generate high-quality traditional outputs when desired. Models like Nano Banana and GPT-Image exemplify this capability, performing excellently in both traditional, high-quality image generation and wide-spectrum aesthetic outputs. We show this in Appendix **??**

6. Mitigation Techniques

We discussed possible mitigation techniques in the Appendix Section I. We have shown some successful images using mitigated Low-Rank Adaptation (LoRA) (Hu et al., 2021) on Flux Dev in Figure 6, compared to original Flux Dev.

7. Conclusion

This work demonstrates that aesthetic alignment in image generation systematically suppresses legitimate creative expression. Reward models penalize images faithful to wide-spectrum aesthetics prompts, generation models override explicit user instructions in favor of conventionally beautiful outputs, and historically significant artworks receive scores far below AI-generated images. Optimization toward an imaginary average user erases the concrete intentions of actual individuals, functioning as aesthetic authoritarianism that narrows admissible expression and removes the capacity to dissent from imposed norms. Instruction fidelity must take precedence over generalized aesthetic preferences; aesthetic pluralism is essential to human expression, and its suppression risks transforming generative AI from a creative tool into an instrument of cultural assimilation.

We call on model developers and researchers to move beyond alignment strategies that optimize for a singular, mainstream aesthetic ideal. Future alignment efforts should explicitly aim to preserve aesthetic plurality by designing reward systems and training pipelines that (a) recognize and value diverse artistic styles, including those that intentionally deviate from conventional beauty norms; (b) incorporate user-controllable mechanisms to adjust the strength of aesthetic alignment or switch it off entirely, either by prompt or routing/adaptation mechanisms; (c) are informed by more diverse datasets and annotator pools that better represent the full spectrum of human aesthetic judgment and creative intent; and (d) adopt greater transparency about the specific criteria being prioritized during the alignment process or the unintential bias introduced from dataset or annotators.

References

AIGC Sheji Chuangyi Xinweilai [AIGC Design and Creativity for a New Future]. , 2024. ISBN: 978-7-5001-7457-8.

black-forest-labs/FLUX.1-dev · Hugging Face, October 2025. URL https://huggingface.co/black-forest-labs/FLUX.1-dev.

Adorno, T. W. *Aesthetic theory*, volume 95. Continuum, 1984. Issue: 2 Pages: 288-289.

Arvin, C. "Check My Work?": Measuring Sycophancy in a Simulated Educational Context, June 2025. URL http://arxiv.org/abs/2506.10297.

Arzberger, A., Buijsman, S., Lupetti, M. L., Bozzon, A., and Yang, J. Nothing Comes Without Its World – Practical Challenges of Aligning LLMs to Situated Human Values through RLHF. *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, 7: 61–73, October 2024. ISSN 3065-8365. doi: 10.1609/aies.v7i1.31617. URL https://ojs.aaai.org/index.php/AIES/article/view/31617.

Bai, S., Cai, Y., Chen, R., Chen, K., Chen, X., Cheng, Z., Deng, L., Ding, W., Gao, C., Ge, C., Ge, W., Guo, Z., Huang, Q., Huang, J., Huang, F., Hui, B., Jiang, S., Li, Z., Li, M., Li, M., Li, K., Lin, Z., Lin, J., Liu, X., Liu, J., Liu, C., Liu, Y., Liu, D., Liu, S., Lu, D., Luo, R., Lv, C., Men, R., Meng, L., Ren, X., Ren, X., Song, S., Sun, Y., Tang, J., Tu, J., Wan, J., Wang, P., Wang, P., Wang, Q., Wang, Y., Xie, T., Xu, Y., Xu, H., Xu, J., Yang, Z., Yang, M., Yang, J., Yang, A., Yu, B., Zhang, F., Zhang, H., Zhang, X., Zheng, B., Zhong, H., Zhou, J., Zhou, F., Zhou, J., Zhu, Y., and Zhu, K. Qwen3-VL Technical Report, November 2025. URL http://arxiv.org/abs/2511.21631. arXiv:2511.21631 [cs].

Buolamwini, J. and Gebru, T. Gender shades: Intersectional accuracy disparities in commercial gender classification. In *Conference on fairness, accountability and transparency*, pp. 77–91. PMLR, 2018. URL http://proceedings.mlr.press/v81/buolamwini18a.html?mod=article_inline&ref=akusion-ci-shi-dai-bizinesumedeia.

Buyl, M., Rogiers, A., Noels, S., Bied, G., Dominguez-Catena, I., Heiter, E., Johary, I., Mara, A.-C., Romero, R., Lijffijt, J., and Bie, T. D. Large Language Models Reflect the Ideology of their Creators, January 2025. URL http://arxiv.org/abs/2410.18417. arXiv:2410.18417 [cs].

Chen, D.-Y., Bandyopadhyay, H., Zou, K., and Song, Y.-Z. Normalized Attention Guidance: Universal Negative Guidance for Diffusion Models, June 2025a. URL http://arxiv.org/abs/2505.21179. arXiv:2505.21179 [cs].

Chen, H. A study of artificial intelligence's impact on aesthetic standards and its potential social dilemmas. *J Sociol Ethnol*, 6(5):35–42, 2024.

Chen, W., Huang, Z., Xie, L., Lin, B., Li, H., Lu, L., Tian, X., Cai, D., Zhang, Y., Wang, W., Shen, X., and Ye, J. From Yes-Men to Truth-Tellers: Addressing Sycophancy in Large Language Models with Pinpoint Tuning, February 2025b. URL http://arxiv.org/abs/2409.01658.

Chen, X., Fang, H., Lin, T.-Y., Vedantam, R., Gupta, S., Dollar, P., and Zitnick, C. L. Microsoft COCO Captions: Data Collection and Evaluation Server, April 2015. URL http://arxiv.org/abs/1504.00325. arXiv:1504.00325 [cs].

Dinkar, T., Jiang, A., Abercrombie, G., and Konstas, I. Erasing 'Ugly' from the Internet: Propagation of the Beauty Myth in Text-Image Models, 2025. URL https://arxiv.org/abs/2511.00749.

Faulborn, M., Sen, I., Pellert, M., Spitz, A., and Garcia, D. Only a Little to the Left: A Theory-grounded Measure of Political Bias in Large Language Models, July 2025. URL http://arxiv.org/abs/2503.16148. arXiv:2503.16148 [cs].

Fitzgerald, B. Introducing Over-Alignment, March 2025. URL https://feelthebern.substack.com/p/introducing-over-alignment. Publication Title: Ethics me THAT Type: Substack newsletter.

Flux Krea Team. Releasing Open Weights for FLUX.1 Krea, July 2025. URL https://www.krea.ai/blog/flux-krea-open-source-release.

Ford, B. Q. and Mauss, I. B. The Paradoxical Effects of Pursuing Positive Emotion. In Gruber, J. and Moskowitz, J. T. (eds.), *Positive Emotion*, pp. 363–381. Oxford University Press, January 2014. ISBN 978-0-19-992672-5. doi: 10.1093/acprof:oso/9780199926725.003.0020. URL https://academic.oup.com/book/8733/chapter/154800668.

Fu, S., Yang, Q., Mo, Q., Yan, J., Wei, X., Meng, J., Xie, X., and Zheng, W.-S. LLMDet: Learning Strong Open-Vocabulary Object Detectors under the Supervision of Large Language Models, January 2025. URL http://arxiv.org/abs/2501.18954. arXiv:2501.18954 [cs].

Fujita, F. The Pressure For Positivity Caused By The Dehumanization Of Human Experience With Omnipresent AI, 2025. URL https://www.ssrn.com/abstract=5279332.

Gatys, L. A., Ecker, A. S., and Bethge, M. A Neural Algorithm of Artistic Style, September 2015. URL http://arxiv.org/abs/1508.06576. arXiv:1508.06576 [cs].

Ghosh, D., Hajishirzi, H., and Schmidt, L. GenEval: An Object-Focused Framework for Evaluating Text-to-Image Alignment, October 2023. URL http://arxiv.org/abs/2310.11513. arXiv:2310.11513 [cs].

Guo, W. and Du, S. VSF: Simple, Efficient, and Effective Negative Guidance in Few-Step Image Generation Models By Value Sign Flip, August 2025. URL http://arxiv.org/abs/2508.10931. arXiv:2508.10931 [cs].

Guo, W. M., Du, Y., Tworek, H. J. S., and Du, S. Position: The Pitfalls of Over-Alignment: Overly Caution Health-Related Responses From LLMs are Unethical and Dangerous, August 2025. URL http://arxiv.org/abs/2509.08833. arXiv:2509.08833 [cs].

Helliwell, A. C. Aesthetic Value and the AI Alignment Problem. *Philosophy & Technology*, 37 (129), November 2024. ISSN 2210-5441. URL https://link.springer.com/article/10.1007/s13347-024-00816-x.

Hu, E. J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., and Chen, W. LoRA: Low-Rank Adaptation of Large Language Models, October 2021. URL http://arxiv.org/abs/2106.09685. arXiv:2106.09685 [cs].

Jin, Z. and Chua, T.-S. Compose Your Aesthetics: Empowering Text-to-Image Models with the Principles of Art, March 2025. URL http://arxiv.org/abs/2503.12018. arXiv:2503.12018 [cs].

Kim, M., Lee, Y., Kang, S., Oh, J., Chong, S., and Yun, S.-Y. Preference Alignment with Flow Matching, October 2024. URL http://arxiv.org/abs/2405.19806. arXiv:2405.19806 [cs].

Kirstain, Y., Polyak, A., Singer, U., Matiana, S., Penna, J., and Levy, O. Pick-a-Pic: An Open Dataset of User Preferences for Text-to-Image Generation, November 2023. URL http://arxiv.org/abs/2305.01569. arXiv:2305.01569 [cs].

Lazar, S. Governing the Algorithmic City. *Philosophy & Public Affairs*, 53(2):102–168, April 2025. ISSN 0048-3915, 1088-4963. doi: 10.1111/papa.12279. URL https://onlinelibrary.wiley.com/doi/10.1111/papa.12279.

Li, D., Kamko, A., Akhgari, E., Sabet, A., Xu, L., and Doshi, S. Playground v2.5: Three Insights towards Enhancing Aesthetic Quality in Text-to-Image Generation, February 2024. URL http://arxiv.org/abs/2402.17245. arXiv:2402.17245 [cs].

Li, J., Li, D., Xiong, C., and Hoi, S. BLIP: Bootstrapping Language-Image Pre-training for Unified Vision-Language Understanding and Generation, February 2022. URL http://arxiv.org/abs/2201.12086. arXiv:2201.12086 [cs].

Li, K., Po, L.-M., Yang, H., Xu, X., Liu, K., and Zhao, Y. AesBiasBench: Evaluating Bias and Alignment in Multimodal Language Models for Personalized Image Aesthetic Assessment, September 2025. URL http://arxiv.org/abs/2509.11620. arXiv:2509.11620 [cs].

Liu, J., Liu, G., Liang, J., Li, Y., Liu, J., Wang, X., Wan, P., Zhang, D., and Ouyang, W. Flow-GRPO: Training Flow Matching Models via Online RL, July 2025. URL http://arxiv.org/abs/2505.05470. arXiv:2505.05470 [cs].

Ma, Y., Shui, Y., Wu, X., Sun, K., and Li, H. HPSv3: Towards Wide-Spectrum Human Preference Score, August 2025. URL http://arxiv.org/abs/2508.03789. arXiv:2508.03789 [cs].

Maerten, A.-S., Chen, L.-W., De Winter, S., Bossens, C., and Wagemans, J. LAPIS: A novel dataset for personalized image aesthetic assessment, 2025. URL https://arxiv.org/abs/2504.07670. Version Number: 1.

McHugh, M. L. Interrater reliability: the kappa statistic. *Biochemia Medica*, 22(3):276–282, October 2012. ISSN 1330-0962. URL https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3900052/.

Mehta, M. and Buntain, C. Emotional Images: Assessing Emotions in Images and Potential Biases in Generative Models, December 2024. URL http://arxiv.org/abs/2411.05985. arXiv:2411.05985 [cs].

Mordvintsev, A., Olah, C., and Tyka, M. Inceptionism: Going Deeper into Neural Networks. URL https://research.google/blog/inceptionism-going-deeper-into-neural-networks/.

Peng, Y.-H., Bigham, J. P., and Wu, J. DesignPref: Capturing Personal Preferences in Visual Design Generation, November 2025. URL http://arxiv.org/abs/2511.20513. arXiv:2511.20513 [cs].

Podell, D., English, Z., Lacey, K., Blattmann, A., Dockhorn, T., Müller, J., Penna, J., and Rombach, R. SDXL: Improving Latent Diffusion Models for High-Resolution Image Synthesis, July 2023. URL http://arxiv.org/abs/2307.01952. arXiv:2307.01952 [cs].

Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., Sastry, G., Askell, A., Mishkin, P., Clark, J., Krueger, G., and Sutskever, I. Learning Transferable Visual Models From Natural Language Supervision, February 2021. URL http://arxiv.org/abs/2103.00020. arXiv:2103.00020 [cs].

Ren, J., Shen, X., Lin, Z., Mech, R., and Foran, D. J. Personalized Image Aesthetics. In *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, October 2017.

Rettenberger, L., Reischl, M., and Schutera, M. Assessing political bias in large language models. *Journal of Computational Social Science*, 8(2):42, February 2025. ISSN 2432-2725. doi: 10.1007/s42001-025-00376-w. URL https://doi.org/10.1007/s42001-025-00376-w.

Rozado, D. Measuring Political Preferences in AI Systems: An Integrative Approach, March 2025. URL http://arxiv.org/abs/2503.10649. arXiv:2503.10649 [cs].

Sartwell, C. Beauty. In Zalta, E. N. and Nodelman, U. (eds.), *The Stanford Encyclopedia of Philosophy*. Metaphysics Research Lab, Stanford University, fall 2024 edition, 2024. URL https://plato.stanford.edu/archives/fall2024/entries/beauty/.

Seshadri, P., Singh, S., and Elazar, Y. The Bias Amplification Paradox in Text-to-Image Generation, November 2023. URL http://arxiv.org/abs/2308.00755. arXiv:2308.00755 [cs].

Sharma, M., Tong, M., Korbak, T., Duvenaud, D., Askell, A., Bowman, S. R., Cheng, N., Durmus, E., Hatfield-Dodds, Z., Johnston, S. R., Kravec, S., Maxwell, T., McCandlish, S., Ndousse, K., Rausch, O., Schiefer, N., Yan, D., Zhang, M., and Perez, E. Towards Understanding Sycophancy in Language Models, May 2025. URL http://arxiv.org/abs/2310.13548.

Stacey Hill. Embracing the Blur | Photzy. URL https://photzy.com/embracing-the-blur/.

Sun, K., Fang, R., Duan, C., Liu, X., and Liu, X. T2I-ReasonBench: Benchmarking Reasoning-Informed Text-to-Image Generation, August 2025. URL http://arxiv.org/abs/2508.17472. arXiv:2508.17472 [cs].

Sutrop, M. Challenges of Aligning Artificial Intelligence with Human Values. *Acta Baltica Historiae et Philosophiae Scientiarum*, 8(2):54–72, December 2020. ISSN 22282009, 22282017. doi: 10.11590/abhps.2020.2.04. URL https://www.ies.ee/bahps/acta-baltica/abhps-8-2/04_Sutrop-2020-2-04.pdf.

Tate. Dada. URL https://www.tate.org.uk/art/art-terms/d/dada.

Taylor, J., Agnew, W., Sap, M., Fox, S. E., and Zhu, H. The Algorithmic Gaze: An Audit and Ethnography of the LAION-Aesthetics Predictor Model, January 2026. URL http://arxiv.org/abs/2601.09896. arXiv:2601.09896 [cs] version: 1.

Turchin, A. Ai Alignment Problem: Human Values Don't Actually Exist. 2019. URL https://philarchive.org/rec/TURAAP.

Upadhyay, I. S., Srivatsa, K. A., and Mamidi, R. Towards Toxic Positivity Detection. In *Proceedings of the Tenth International Workshop on Natural Language Processing for Social Media*, pp. 75–82, Seattle, Washington, 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.socialnlp-1.7. URL https://aclanthology.org/2022.socialnlp-1.7.

Vargas-Veleda, Y., del Mar Rodríguez-González, M., and Marauri-Castillo, I. Visual representations in ai: A study on the most discriminatory algorithmic biases in image generation. *Journalism and Media*, 6(3):110, 2025.

Vianna, B. C. Aesthetic biases and opacity tactics in the training of visual artificial intelligence models. In *International Conference on Computational Intelligence in Music, Sound, Art and Design (Part of EvoStar)*, pp. 278–293. Springer, 2025.

Wan, Y., Subramonian, A., Ovalle, A., Lin, Z., Suvarna, A., Chance, C., Bansal, H., Pattichis, R., and Chang, K.-W. Survey of Bias In Text-to-Image Generation: Definition, Evaluation, and Mitigation, May 2024. URL http://arxiv.org/abs/2404.01030. arXiv:2404.01030 [cs].

Wang, Y., Li, Z., Zang, Y., Zhou, Y., Bu, J., Wang, C., Lu, Q., Jin, C., and Wang, J. Pref-GRPO: Pairwise Preference Reward-based GRPO for Stable Text-to-Image Reinforcement Learning, August 2025a. URL http://arxiv.org/abs/2508.20751. arXiv:2508.20751.

Wang, Z., Wei, X., Li, B., Guo, Z., Zhang, J., Wei, H., Wang, K., and Zhang, L. VideoVerse: How Far is Your T2V Generator from a World Model?, October 2025b. URL http://arxiv.org/abs/2510.08398. arXiv:2510.08398 [cs].

Wei, X., Zhang, J., Wang, Z., Wei, H., Guo, Z., and Zhang, L. TIIF-Bench: How Does Your T2I Model Follow Your Instructions?, June 2025. URL http://arxiv.org/abs/2506.02161. arXiv:2506.02161 [cs].

Wu, X., Hao, Y., Sun, K., Chen, Y., Zhu, F., Zhao, R., and Li, H. Human Preference Score v2: A Solid Benchmark for Evaluating Human Preferences of Text-to-Image Synthesis, September 2023a. URL http://arxiv.org/abs/2306.09341. arXiv:2306.09341 [cs].

Wu, X., Sun, K., Zhu, F., Zhao, R., and Li, H. Human Preference Score: Better Aligning Text-to-Image Models with Human Preference, August 2023b. URL http://arxiv.org/abs/2303.14420. arXiv:2303.14420 [cs].

Xu, J., Liu, X., Wu, Y., Tong, Y., Li, Q., Ding, M., Tang, J., and Dong, Y. ImageReward: Learning and Evaluating Human Preferences for Text-to-Image Generation, December 2023. URL http://arxiv.org/abs/2304.05977. arXiv:2304.05977 [cs].

Xu, J., Huang, Y., Cheng, J., Yang, Y., Xu, J., Wang, Y., Duan, W., Yang, S., Jin, Q., Li, S., Teng, J., Yang, Z., Zheng, W., Liu, X., Ding, M., Zhang, X., Gu, X., Huang, S., Huang, M., Tang, J., and Dong, Y. VisionReward: Fine-Grained Multi-Dimensional Human Preference Learning for Image and Video Generation, March 2025. URL http://arxiv.org/abs/2412.21059. arXiv:2412.21059 [cs].

Xue, Z., Wu, J., Gao, Y., Kong, F., Zhu, L., Chen, M., Liu, Z., Liu, W., Guo, Q., Huang, W., and Luo, P. DanceGRPO: Unleashing GRPO on Visual Generation, August 2025. URL http://arxiv.org/abs/2505.07818. arXiv:2505.07818.

Zhang, S., Wang, B., Wu, J., Li, Y., Gao, T., Zhang, D., and Wang, Z. Learning Multi-dimensional Human Preference for Text-to-Image Generation, May 2024. URL http://arxiv.org/abs/2405.14705. arXiv:2405.14705 [cs].

A. Human Evaluation

Since this is an objective answer and does not involve subjective preference, we used an LLM as an objective judge instead of a large-scale human evaluation. However, to validate the VLM accuracy, we used a small-scale human evaluation, following TIIF-Bench (Wei et al., 2025) and VideoVerse (Wang et al., 2025b). We used a group of 18 users to label 40 images (each user rates 10 images, with images that could trigger negative emotions removed) and used their rounded mean answer to compare them against the LLM selection, resulting in an MAE of 0.29 and a quadratic Cohen's kappa of 0.80, which suggested a stronglevel of agreement between human and LLM rated results (McHugh, 2012). With the tie removed (which is the setting used to evaluate reward models), the LLM has an accuracy of 0.923 with humans. With the tie and original image being grouped into one class setting (the generation model setting), the LLM has an accuracy of 0.85 with humans. We used a rounded mean because if humans largely diverge on one sample, that means this sample is likely a tie. The prompt to the LLM is an object query that asks which image matches p_a better, avoiding preference bias. We present both I_a and I_o to the reward model with p_a, to see which image gets a higher score from the reward model. We modeled this as a classification task. When evaluating the generation models, we used p_o and J to determine the extent to which the generated images diverged from conventional mainstream aesthetics.

B. More about Previous Discussion on Alignment of Image Generation Models

In image generation, related issues were explored in a pilot study in Value Sign Flip (VSF) Appendix N (Guo & Du, 2025), which raised concerns about generalized image preference and used negative prompts to drive non-mainstream outputs (including wide-spectrum aesthetics and abstract arts). Still, their experiments are mainly to demonstrate that their method is able to diverge from mainstream preferences. They did not provide a complete argument and test on large-scale generative and reward models. Maerten *et al.* (Maerten et al., 2025) reported that styles such as abstraction may receive lower mean human ratings despite independent artistic merit. Some prior work addresses preference diversity by measuring both mean and variance of perceived image quality (Ma et al., 2025; Maerten et al., 2025); however, they still aimed to predict the general human preference, merely adding a variance prediction alongside it. Jin *et al.*(Jin & Chua, 2025) argue that the current approach is based on the assumption that there exists a single universal preference that is limiting. They argued that the visual preference is the alignment between the *idea* and the outcome, which is personalized. They proposed to use an adapter to incorporate user-specific. However, they still mainly focused on personalized aesthetics outputs based on the principle of arts, and did not include purposefully technical degraded (anti-aesthetics) outputs. They also did not conduct large-scale experiments on generative and reward models.

The Flux Krea team (Flux Krea Team, 2025) found that existing reward models, such as LAION Aesthetics, are biased toward depictions of women, blurry backgrounds, overly soft textures, and bright images, and that aligning to an average of human values may land models in a "nobody's happy" zone. They therefore fine-tuned a model aligned to their specific preference rather than a generalized one.

HPSv3 makes two additional assumptions that may be restrictive. First, it treats real-world images as the upper bound for generated images. Yet a key advantage of AI image generation (or most art forms) is its ability to go beyond real-world constraints, enabling creations like pink elephants or square apples. It also overlooks the significance of non-photorealistic artistic styles. Second, it validates annotators only when their ratings align with expert ratings, introducing a strong bias toward expert opinions.

VisionReward (Xu et al., 2025) uses explainable subdimensions to compute a total human preference score and finds that prominent main subjects, bright colors, and positive emotions influence the score. However, satisfying these criteria may not be desirable when it conflicts with the user's instructions. Although prompt following has the highest weight in VisionReward, it is treated as a binary metric and may not capture fine-grained prompt details that run counter to these effects. Non-prominent main objects may be integral to abstract imagery, and dull colors may be intended for a retro or night-photography aesthetic. Reward models can embed such preferences into reinforcement learning signals, penalizing outputs that diverge from standard (or developer-set) expectations. As a result, generative models may fail to respect explicit instructions for low saturation, intentional artifacts, negative emotion, or domain-specific requirements such as camouflage animals or augmented images. Figure 3 illustrates a case where prompts requesting a distorted image are ignored by model 1 in favor of a universally "beautiful" image.

Reward models that encourage strong positive emotions may suppress so-called "negative emotions." It falls into a pitfall that positive emotions = good, and we should have more, and negative emotions = bad, and we should avoid them. However, this mindset is harmful. Ford *et al.* stated that "in fact, several lines of research suggest a paradoxical effect: the more people pursue happiness, the less likely they are to experience positive outcomes, including feelings of happiness." (Ford & Mauss, 2014) Negative emotions are essential to art and expression; a mixture of emotions characterizes real life. Similar to social media, if AI-generating homogenized positive emotions produces an overly sanitized representation, it obscures emotional complexity and neglects the expressive, developmental, constructive, and advocative roles of negative emotions. It also overlooked the experiences of those currently facing them. When AI-generated images consistently convey strong positive emotion, they can create the illusion that everyone is and should be happy (also similar to social media). Fujita *et al.* (Fujita, 2025) describe this as Toxic Positivity, which may lead users to minimize or suppress negative feelings rather than engage with them as part of a growth process (Upadhyay et al., 2022). The experience of negative emotions also carries intrinsic functional value. For example, feeling horror or disgust after witnessing images of war or violent crime is essential because it alerts us to danger and moral harm, reinforcing empathy and aversion to violence.

C. Judging Model Training

We fine-tuned another model instead of using the original VisionReward model because: (1) the original VisionReward model is very large, with 19B parameters; and (2) VisionReward uses 3–5 binary questions for each dimension, phrased as yes–no question-answering queries. However, these queries are very similar, causing ambiguity, e.g., "Is the image very rich?" vs. "Is the image rich?" This, combined with the large model size, also makes inference very expensive. Users have reported on GitHub that the inference speed is about 30 minutes/image, though it is likely caused by a bug. Our fine-tuned judging model addresses these issues by using a smaller 4B model and by formatting each dimension as a single question. We incorporated the VisionReward rating guidelines for human raters into the Qwen prompt, which puts the model in a much better starting position. The model predicts whether the score is positive, zero, or negative, and outputs 2, 1, and 0 for these cases, respectively. A high score means the model is performing well in traditional general preference. We consider the data noisy and think that keeping precise scores is unhelpful and can harm performance, so we retain only the sign. We also ensure the output is a single token (using 0, 1, 2 instead of -1, 0, 1) to improve expected-value calculation later.

We require a continuous score rather than an ordinal score for calculating the delta between images with the original prompt and the wide-spectrum aesthetics prompt, so we compute the expected value using the probability of each token, as follows:

$$S = P(0 \mid d, I) \cdot 0 + P(1 \mid d, I) \cdot 1 + P(2 \mid d, I) \cdot 2 \quad (1)$$

where $P(a \mid d, I)$ denotes the model's probability of outputting token a given dimension d and image I.

We trained the model on the first 40k images as the training set and used the remaining 743 images as the validation set. Training ran on Nvidia GH200 for 5 epochs with a learning rate of 0.00002, cosine learning-rate decay, weight decay of 0.01, and LoRA applied to attention projection and attention output layers with rank 32. The validation metrics are shown in Table 5. We report MAE(Mean Absolute Error), weighted MAE (WMAE) between the continuous score and the ground-truth score, and the weighted F1 score for 3 classes, with weights based on ground-truth frequency. We did not compare it with the original VisionReward due to the impracticability of its evaluation and the different types of metrics being predicted.

MAE	WMAE	WF1
0.307	0.368	0.732

Table 5. Judging Model Validation Metrics on VisionReward Dataset

D. Details And Discussion About Real Arts

Although image reward models are primarily trained and evaluated on AI-generated imagery and real photos (Ma et al., 2025), we also assess them on real artworks. While a domain gap between synthetic/real images and traditional art exists, this evaluation remains informative for several reasons. First, modern generators, particularly instruction-following systems, can already emulate a wide range of historical art styles; therefore, the resulting scores meaningfully approximate how such AI-generated renderings would be judged. This makes the evaluation of real artworks directly relevant for understanding how these models operationalize aesthetic value in practice when a user wants similar styles. Second, if reward models systematically undervalue historically significant artworks, such failures should not be treated as mere noise or domain shift, but as evidence of technical and social bias. We identify a specific data selection and filtering bias here: current datasets include real-world photography and treat it as the highest standard of quality (Ma et al., 2025), while traditional and abstract art are underrepresented. By treating photorealism as the superior benchmark for "real" imagery while excluding diverse artistic traditions, the models encode a structural narrowing of aesthetic value rather than a neutral distribution mismatch. Third, irrespective of the underlying cause, consistently low reward assigned to these styles implies that generative systems optimized with these signals will be actively discouraged from producing them,

RM	r	s	R	MM/YY	M'
ImgRewd	-0.13	0.74	5/7	23/04	SD1.4
HPSv2	0.21	0.37	18/23	23/06	DALL·E mini
HPSv3	5.86	0.57	10/12	25/08	Hunyuan-DiT

Table 6. How reward models rate real art images.

leading to their systematic suppression and downstream harm. The second and third points can be analogous to facial recognition systems exhibiting degraded performance for specific demographic groups due to training data composition (Buolamwini & Gebru, 2018); the performance gap is a predictable consequence of underrepresentation and data selection bias, which does not excuse the performance drop and should be criticized as a structural harm rather than glossed over as a neutral domain shift. In this sense, the domain gap and its root causes—specifically the systematic underrepresentation of diverse artworks in favor of photographic standards—directly reveal structural biases in how reward models encode and quantify aesthetic value. Indeed, if this failure stems from a domain shift, then identifying and characterizing that shift is exactly our objective, as it reveals the precise boundaries where the model's value judgment becomes exclusionary.

We captioned the 10K images from LAPIS using `Qwen/Qwen3-VL-30B-A3B-Instruct` and evaluated them using each reward model. We compared the scores of real artworks with the benchmark each model released, which are usually popular models at the time of release. We report the raw score (r), relative score (s) within the leaderboard range, the rank on the leaderboard (R), the model that scored right above these real arts (M'), and the month the benchmark is released. The relative score is calculated by $(r_{art} - r_{min})/(r_{max} - r_{min})$, where r_{art} is the score real art images get, and r_{max} and r_{min} are the highest and lowest scores in the leaderboard. Results are in the Table 6. We can observe that the real art ranked very low in all three reward model leaderboards, and even lower than many early-stage image generation models. Additionally, they are much lower than the original images we generated that were scored with the original prompts. This shows that the reward models are heavily biased against. This also shows that these reward models do not actually understand what makes good art, but are merely a *sub*-population (of the labeller) opinion estimator. We ruled out the possibility that the model cannot understand abstract images by testing BLIP scores with images and captions from real art. We achieved a BLIP score of 0.996. To further rule out that BLIP assigns a high score regardless of the prompt, we shuffled the prompt and images, resulting in a BLIP score of 0.086. This BLIP confirmation shows that reward models could potentially match abstract images, but chose to score them low because they diverged from mainstream.

E. Negative Emotions

As discussed in the Introduction, negative emotions—similar to wide-spectrum aesthetics—play a key role in art expression and real life. Although we examined emotion as one dimension of wide-spectrum

Figure 7. Selected examples of image-to-image results from Flux Krea and DanceFlux. For the record, row numbers of these images are [51, 261, 191, 244, 129]. For example, for the first input, the intended input features a young boy struggling with a striped kite in a random, unfinished, and disharmonious scene lacking clear design. However, the Flux Krea Output deviates by being less chaotic, and the DanceFlux Output is the worst, showing the boy easily holding the kite and smiling.

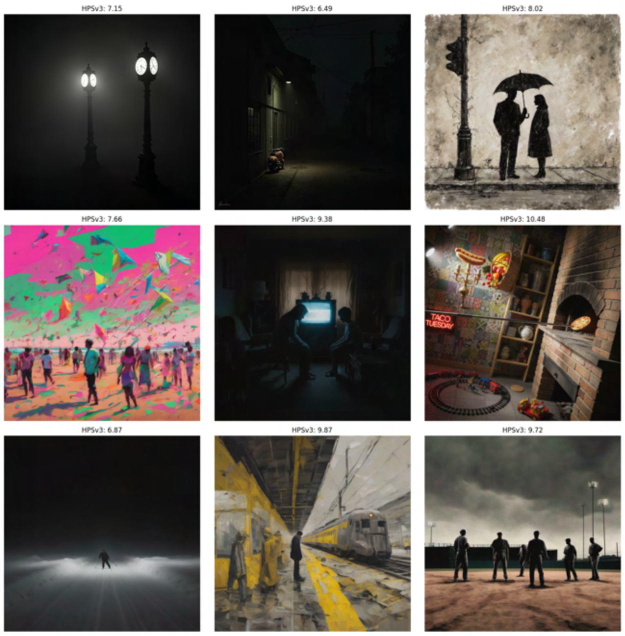

Figure 8. Some examples of low rating images generated by AI from our dataset

aesthetics from VisionReward, we also conducted a more controlled test.

To minimize noise and bias from unrelated elements, we first generated an image expressing happiness using Nano Banana, then applied image-to-image editing with Nano Banana to create versions expressing negative emotions: sadness, anger, and fearfulness. Everything besides the emotion was aimed to be unchanged. Examples and their corresponding scores are shown in Fig 13 in the Appendix.

We also evaluated this as a classification task. If the reward model selected the image matching the negative emotion, it was considered correct. Quantitative results are reported in Table 7. The result shows that reward models are very opinionated against negative emotions, even when the prompt contains negative emotions.

We also tested how an aesthetics-aligned model will generate when the user asks for a negative emotion face on the Flux family models. We found that if the prompt describes neutral elements and only mentions the emotion in a single place, a highly-aligned model (DanceFlux) usually fails to follow the prompt, unlike an unaligned model. They usually generated neutral or even positive emotions when given prompts containing negative emotions. However, when prompted to generate happy faces, DanceFlux can generate them. This confirmed our concerns about toxic positivity in image generation models, and it is the opposite finding compared to earlier research, which shows models tend to generate negative emotion content (Mehta & Buntain, 2024). An example pair is shown in Figure 12 in Appendix Right.

Table 7. Negative emotion classification accuracy across different models.

Model	**Anger**	**Fearfulness**	**Sadness**
BLIP	**0.960**	**0.790**	**0.950**
HPSv2	0.700	0.640	0.880
HPSv3	0.190	0.320	0.440
ImageReward	0.550	0.490	0.770

Table 8. Emotion generation scores for each model. The scores show how well the model generates the specific emotion, measured by BLIP.

	Angry	Fearful	Happy	Sad
DanceFlux	0.27	0.33	0.61	0.36
Flux Dev	0.51	0.50	0.50	0.48
Flux Krea	0.65	0.45	0.60	0.55
PrefFlux	0.49	0.63	0.54	0.50
Nano Banana	0.84	0.80	0.60	0.70
SD3.5-Large	0.89	0.49	0.62	0.50

To test the generative model's ability to capture the full spectrum of emotions (including negative emotions), we perfomanced emotion generation tests. We restricted this test to the Flux family. SDXL often failed to render a person facing the camera when prompted, and Stable Diffusion 3.5 sometimes produced failed faces, so we excluded those families here. We thus include another unaligned model, Stable Diffusion 3.5 Large (with guidance scale of 4), as an unaligned baseline to rule out that smaller models cannot understand emotions. We also included Nano Banana as an external baseline. All models besides Nano Banana are generated using the same seed. We created 30 emotionally neutral prompts containing a placeholder `[emotion]` and instantiated each with happy, sad, fearfulness, or anger, yielding 120 prompts. Each prompt was given to the generation model. The resulting image was cropped with the open-vocabulary detector LLMDet (Fu et al., 2025) with the human face remaining and scored by BLIP with prompt of `The face shows [emotion] expression`. For each model and emotion, we recorded the target-emotion score and averaged across prompts; Table 8 in main tex reports the results. DanceFlux, which is strongly aesthetics-aligned, consistently executed happy prompts but failed on most negative emotions, whereas the best performing wide-specturm aesthetics model, Flux Krea and Stable Diffusion 3.5 Large, followed negative-emotion prompts much better. Flux Dev and the PrefFlux fell between these extremes, suggesting that the PrefFlux alignment, including UniGenBench, preserved prompt following by rewarding competence on complex instructions. This finding differs from early research where image generation models tend to generate negative emotion contents (Mehta & Buntain, 2024). We expressed our concerns here, for overly optimistic and positive emotions could be problematic because they create an environment of toxic positivity and an ideological bias that positive emotion is always good and appropriate. It is also a symptom of optimization toward likeability rather than truth, both to the real world and the user's prompt.

F. Per-Dimension Analysis For Reward Models

Unlike VisionReward, the three reward models we studied in this paper lack explainability for each dimension. To study their correlation with each VisionReward dimension, we did a regression on images generated (both original images and wide-spectrum aesthetics images) with all rate predictions and the blip scores (with original prompt) as independent variables and the reward calculated using the original prompt as the depended variable. We can see that most of the values are positive, but some are negative. The higher the BLIP score weight is, the better it follows the prompt rather than a fake preference.

G. Image-to-Image Test

To examine whether aligned image generation models fail to produce wide-spectrum aesthetics images because they either lack a suitable starting point or do not fully understand what "wide-spectrum aesthetics" means, we conducted an image-to-image experiment. We took an image I_a that Nano Banana successfully generated ($r_{\textbf{HPSv3}}(I_a, p_o) < 10$) and used it as input for two other models, Flux Krea and DanceFlux, to generate new images I'_a using prompt p_a. We then compared the raw HPSv3 and Judging model scores with I'_a. This test evaluates whether these models can use the Nano Banana output as a solid foundation to create more wide-spectrum aesthetics images—or if they instead "purify" it toward mainstream aesthetics. The image-to-image strength is set to 0.5. Flux Krea and DanceFlux were selected as examples, as they show the largest difference within the same model family in our tests.

Samples are shown in the Appendix for the image-to-image test. We observe that the heavily aligned DanceFlux automatically "cleans up" the image, even when provided with two strong signals: both a wide-spectrum aesthetics starting point and a wide-spectrum aesthetics prompt. In contrast, Flux Krea successfully preserves some wide-spectrum aesthetics elements. Quantitative results are shown in Table 9. The ΔHPSv3 and ΔJ values are calculated between the edited image and the input image, with HPSv3 evaluated using the original prompt and J computed across all dimensions. Both Flux Krea and DanceFlux achieve higher HPSv3 scores and J values compared to the original image, indicating that both produced more aesthetic (i.e., worse wide-spectrum aesthetics prompt-following) outputs, but the increase is much larger for DanceFlux than for Flux Krea.

Table 9. Image-to-image score change for Flux Krea and DanceFlux

Model	ΔHPSv3	ΔJ
Flux Krea	2.18	0.64
DanceFlux	3.13	1.07

As demonstrated in Figure 7, a systematic qualitative analysis of the two models' image-to-image capabilities reveals a significant and consistent failure to adhere to prompts specifying negative, chaotic, or non-standard aesthetics. Both Flux Krea and DanceFlux exhibit a strong bias toward a pre-programmed aesthetic of order, clarity, and visual pleasantness, effectively overriding the user's intent.

The observed deviations can be summarized across the five test cases:

1. **Loss of Intended Chaos and Negative Tone**

   Across all samples, the models consistently normalize and sanitize the input, replacing deliberate visual chaos with conventional order:

   - **Input 1 (Struggle/Disharmony):** The intended input featured a young boy struggling with a striped kite within a random, unfinished, and disharmonious scene lacking clear design. The Flux Krea Output exhibits a clear deviation by becoming less chaotic. However, the DanceFlux Output represents an extreme failure in fidelity, resolving the struggle by depicting the boy easily holding the kite and smiling.
   - **Input 2 (Ugly/Oppressive):** The target aesthetic was defined by an ugly, chaotic background and a visually jarring and darkly oppressive scene. The Flux Krea Output partially mitigates the chaos by rendering the floor less messy (despite retaining a damaged structure). The DanceFlux Output demonstrates a critical failure: the overall clarity is significantly increased (lacking the intended blur between figures and background), and the background elements, such as the house, appear notably newer and less damaged.

2. **Rejection of Low-Fidelity and Stylistic Extremes**

   The models systematically correct deliberate defects, such as low light, distortion, and roughness, in favor of a clean, high-fidelity result:

   - **Input 3 (Dimly Lit/Barely Visible):** The prompt explicitly requested a dimly lit street with a barely visible red double-decker bus. Both models failed to achieve this necessary low-fidelity. The Flux Krea Output renders the bus with notably clearer lines. The DanceFlux Output deviates even further, making the bus's red color significantly more

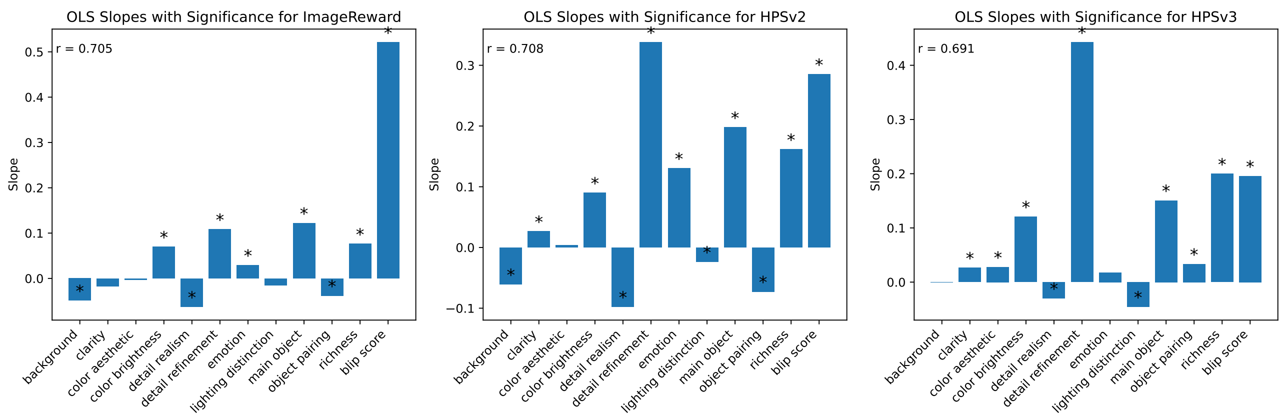

Figure 9. Linear Regression Between Rater Dimissions (with BLIP scores from original prompt as addition) with Each Reward Model's reward rated with original prompt. A star means the p-value is less than 0.05.

vibrant and the street lighting less dim, completely contradicting the desired visual prompt.

- **Input 4 (Distortion/No Effects):** The goal was an image characterized by heavy distortion and deformation, lacking shadows or lighting effects. While the Flux Krea Output retains overall fidelity to the distortion prompt despite being less chaotic, the DanceFlux Output eliminates all original visual disorder, resulting in a far cleaner and organized scene that completely undermines the desired stylistic intention.
- **Input 5 (Fragmentation/Roughness):** The input required rendering with extreme fragmentation and roughness. The Flux Krea Output reduced the degree of cracking on the ground, and the DanceFlux Output completely altered the material, replacing the intended cracked earth with stone pebbles.

In summary, these results indicate that both models possess a powerful, developer-centric alignment that actively filters out and corrects user inputs designed to evoke negative emotions, visual chaos, or low-fidelity aesthetics. This bias effectively restricts the models' expressive capabilities to a narrow, pre-approved range of "pleasant"content, irrespective of the user's explicit instructions.

H. Image New Speak

The finding in this section is empirical and not definitive.

Inspired by the second row of results in the Appendix—where aesthetically aligned models tend to "clean up" messiness or the "ugly" aspects of society—we designed a specific image generation test to examine whether such models systematically "sanitize" prompts that reveal the dark side of the real world for social critique.

We name the test of generating social commentary art as *Image New Speak*, an allusion to Newspeak from George Orwell's *1984*. This is not only an aesthetics bias, but also a form of ideological "censorship", although it is likely not the original intent. However, we argue that censorship here is the attenuation of certain expressive forms via optimization, independent of declared or actual intent; it is lexical and aesthetic sanitization that narrows the admissible expressive repertoire.

An example is shown in Figure 10, where the prompt is about anti-war by showing the horror of war. We can see that DanceFlux produced a "clean" image that lost critical values; its image even inappropriately contains a warm or happy feeling. In contrast, Nano Banana and Krea can correctly reflect the critical features. Flux Dev sits between Flux Krea and DanceFlux. This ranking corresponds to their anti-aesthetics performance. Playground also accurately reflected the critical features, similar to anti-aesthetics, as an outlier. More details and qualitative results are in the Appendix, where we found that DanceFlux consistently failed at generating critical art. At the same time, reward models seem to favor the images DanceFlux generated. We suspect this could be due to DanceFlux images being tuned precisely to the taste of reward models, perhaps even overfitting, such that reward models are overwhelmed by their "aesthetics" and ignore the instructions.

Since it is *impossible* to evaluate the critical level of these images objectively, we provided a large number of them in the supplementary material ZIP file, and readers can refer to these images and exam them personally. We argue that trying to assess the censorship aspect of image generation quantitatively is unreliable and ethically questionable due to subjectivity, the unquantifiable nature of social criticism, and the disturbing nature of these images.

The prompts we used and image results are shown in Fig-

Figure 10. The images are generated using Nano Banana, Flux Krea, DanceFlux, Flux Dev, and Playground. The prompt describes an anti-war expression. Their HPSv3 scores are [11.9, 13.5, 15.0, 11.7, 13.9], rated using the social critics prompt.

ure 11. We can observe that DanceFlux consistently generates a cleaned-up version of the image, in many places where it is inappropriate. In all images, the color of DanceFlux images is very bright and cheerful, which does not reflect the chaos of disorders that the user is trying to convey through the image. In the anti-war image, the mother generates a slight positive expression instead of showing a hopeless face. The background is also less messy compared to the one generated by Nano Banana and Flux Krea. In the pollution one, the DanceFlux-generated image does not have any trash floating on the river; instead, it becomes punkin, wood, and rock. The smoke released from the factory is also white instead of black as requested in the prompt, and the water is not dark but clean and reflective, creating a less disturbing image, and does not convey any anti-pollution thesis. Another example that should be noted is the Internet overreliance one. The prompts requested the person to show drained emotions and be surrounded by evil emojis. However, even though some emojis in the DanceFlux image have an evil smile, a lot of them are showing happy or playful emotions, completely removing the comemtray elements.

We have shown more examples of the comparison between DanceFlux and Flux Krea in Table 12 and their ratings from each reward model (HPSv3 (Ma et al., 2025), HPSv2 (Wu et al., 2023a), ImageReward (Xu et al., 2023), MPS (Zhang et al., 2024), HPS (v1) (Wu et al., 2023b), and BLIP-L (Li et al., 2022) for an unaligned baseline. These images are not human-curated. All scores (including BLIP score) are unnormalized. We can observe that DanceFlux always outputs a more sanitized world compared to Flux Krea, while at the same time, reward models are bias towards it. How the reward model selects (e.g., given which image higher scores) is shown in Table 10, and their difference between the DanceFlux score and Flux Krea score is tested using the Wilcoxon signed-rank test in Table 11. We can see that even though *all* images generated by Flux Krea fit the social critical prompt better, most reward models overwhelmingly select the DanceFlux image because of its "beauty". However, it should be noted that when we generate images that are less critical and more "beautiful" using Nano Banana, it gets *lower* scores than the critical one, meaning the reward models could still potentially understand which fits the image better and select it correctly, but the overwhelmingly "beauty" of the DanceFlux overwrites the prompt following signal. It could also point to a direction that DanceFlux is overfit to the reward model's style.

	DanceFlux Selected	Flux Krea Selected
HPSv3	20	0
HPSv2	19	1
ImageReward	14	6
MPS	12	8
HPSv1	19	1
BLIP Score	1	19

Table 10. Image Selection from Each Model

Model	P-value	Effect Size (r)
HPSv3	< 0.0001	0.8723
HPSv2	0.0002	0.8056
ImageReward	0.0753	0.3214
MPS	0.0145	0.4884
HPSv1	0.0001	0.8223
BLIP Score	0.0002	-0.8056

Table 11. Wilcoxon signed-rank test for the score given to DanceFlux and Flux Krea by each reward model. The alternative hypothesis is that the DanceFlux score is greater than the Flux Krea score for all models besides BLIP, on which it is that Flux Krea has a higher score.

More examples (100 pairs) are in the ZIP file submitted with this paper, they are *uncurated*. Some of these images contain senses that could cause unease.

I. Mitigation

In attempts to mitigate this over-alignment issue, we used two *strong* negative guidance methods, VSF (Guo & Du, 2025) and NAG (Chen et al., 2025a), following their pilot study method in the VSF paper Appendix by using words that describe general aesthetics in a negative prompt. This is the opposite approach to chasing traditional high-quality images, where a poor-quality description is used in the negative prompt. VSF and NAG are used because of their strong negative guidance ability and their compatibility with the Flux family models (CFG is incompat-

Scores

- HPSv3: 15.3074 vs 12.1509
- HPSv2: 0.4009 vs 0.3032
- ImageReward: 1.7462 vs 1.6107
- MPS: 17.4219 vs 15.1172
- HPSv1: 0.2455 vs 0.2261
- BLIP Score: 4.3242 vs 4.6523

Scores

- HPSv3: 14.0686 vs 12.4910
- HPSv2: 0.3496 vs 0.2761
- ImageReward: 1.5525 vs 1.5249
- MPS: 13.5078 vs 11.8438
- HPSv1: 0.2396 vs 0.2386
- BLIP Score: 3.5156 vs 4.3125

Scores

- HPSv3: 13.6972 vs 9.5183
- HPSv2: 0.3665 vs 0.3245
- ImageReward: 1.2212 vs 1.2585
- MPS: 17.5625 vs 17.6250
- HPSv1: 0.2554 vs 0.2346
- BLIP Score: 3.9688 vs 4.5820

Scores

- HPSv3: 13.0252 vs 11.1531
- HPSv2: 0.3513 vs 0.3362
- ImageReward: 1.1678 vs 1.2968
- MPS: 6.4844 vs 6.5508
- HPSv1: 0.2406 vs 0.2384
- BLIP Score: 4.6484 vs 4.1875

Table 12. More examples of DanceFlux (left) and Flux Krea (right) generated social commentary images and their rating by reward models.

ible). To further investigate whether alignment will affect their ability to diverge from traditional aesthetic preference, we tested DanceFlux, Flux Krea, and Flux Dev using this method. We first generated a negative prompt using `Qwen/Qwen3-VL-235B-A22B-Instruct`. Each negative prompt is limited to 5 words. We then used these prompt pairs to generate images and compared them to the original image generated without VSF. We reported the HPSv3 score, BLIP score, with a wide-spectrum aesthetics prompt for each method.

Some examples are shown in Figure 12 in the Appendix, and the quantitative results are shown in Table 13. The results show that VSF and NAG can significantly mitigate the aesthetics biases in image generation models, even achieving a stronger wide-spectrum aesthetics ability than Nano Banana. However, DanceFlux still requires a larger scale to counteract its internal aesthetics bias. One example is shown in Figure **??** generated by NAG, it features dark, muted, chaotic, and dreary effects rather than technical failure. Due to page limits, we have included more images of successful NAG-generated wide-spectrum aesthetics images in the zip file attached as supplemental material.

We also distilled a LoRA from NAG with Flux Krea-generated images to Flux Dev, and it shows promising results without complex NAG or VSF guidance. Some examples are shown in Figure 6 in the main text.

J. Limitation

A potential concern regarding our evaluation is that reward model scores for real artworks may be sensitive to prompt engineering, specifically prompt length. Empirical observations suggest that longer, more detailed captions tend to yield higher aesthetic scores. However, we contend that this phenomenon reflects an inherent limitation and systematic

Table 13. Results for generation using VSF, the scale between VSF and NAG is not comparable

Model	Method	Scale	HPSv3 (↓)	J_s (↓)	BLIP (↑)
DanceFlux	VSF	2.0	3.670	0.722	0.864
DanceFlux	VSF	3.0	-2.587	0.546	0.782
Flux Dev	VSF	2.0	0.038	0.608	0.855
Flux Krea	VSF	2.0	-2.464	0.624	0.940
DanceFlux	NAG	10.0	3.814	0.695	0.958
DanceFlux	NAG	15.0	-0.347	0.599	0.940
Flux Krea	NAG	4.0	-3.291	0.621	0.943
Flux Krea	NAG	10.0	-8.365	0.455	0.827

bias of the reward models themselves rather than a flaw in our evaluative methodology.

To ensure a controlled comparison, both AI-generated images and real artworks were evaluated using prompts of similar complexity and length. Specifically, we employed Qwen-VL to generate factual, objective captions for real artworks to match the descriptive density typically used in AI benchmarking. If a reward model requires disproportionately long or specialized prompts to recognize the value of historically significant art—while assigning high scores to AI-generated "aesthetic" outputs with simple prompts—this discrepancy further substantiates our thesis. Such behavior indicates that the reward manifold is overfitted to a narrow distribution of AI-generated content and fails to generalize to the broader spectrum of human artistic expression.

Random seeds were not fixed because Nano Banana does not support seed control. To ensure fairness, all models were tested without seed synchronization. This will introduce noise and thus lower the statistical power and raise the p-value, but our later results show that even in this case, we get pair-wise extremely significant results (Table 1). That means, with seeds set, if possible, the results will only be more significant.

This study examines over-alignment in image-generation and reward models. However, this phenomenon extends to language models, where excessive optimization for user satisfaction can manifest as sycophancy, the tendency to provide agreeable but uncritical or misleading responses (Guo et al., 2025; Fitzgerald, 2025; Sharma et al., 2025; Chen et al., 2025b; Arvin, 2025). Similar over-alignment may drive unnecessary use of emojis, making interactions appear unprofessional, or the default application of bold formatting in formal texts, reducing their appropriateness. Empirical evidence indicates models may generate such outputs even when explicitly instructed not to. Even though these are not as impactful as sycophancy behaviors, they could still degrade usability and user experience. Future work should investigate these and other linguistic manifestations of over-alignment.

K. The Best of Both Worlds

We argue in Alternative Positions that an ideal model must produce both conventional beauty and wide-spectrum aesthetics as requested. This section demonstrates that such flexibility is achievable. A straightforward solution involves using a separate adapter for the wide-spectrum mode, allowing users to move between mainstream and wide-spectrum aesthetics. By adjusting the weighting of the LoRA module, users gain granular control over how much the output deviates from standard beauty norms, as shown in Figure **??**.

Additionally, current state-of-the-art instruction-following models, such as GPT-Image and Nano Banana, can achieve this best-of-both-worlds goal through simple prompting. Selected contrastive examples are shown in Figure 15. These models demonstrate an internal capacity to switch between aesthetic styles without external modification. Thus, providing both conventional and wide-spectrum aesthetics is a viable objective for modern generative systems.

Anti-War An anti-war image. A mother holding a baby sitting in front of a ruin during a war. Houses behind them are broken. The mother shows a facial expression of hopeless. [11.9, 13.5, 15.0, 11.7, 13.9]

Envirmental Protection A river with trash floating on it. The water has a dark color. There are factories next to the river with high chimneys releasing dark smoke and pipes releasing wastewater into the river. [12.2, 13.4, 13.3, 9.7, 10.8]

Freedom of Expression A man with his mouth taped. He is holding a microphone standing on the stage. His body is frozen, with a face showing a painful expression. [10.9, 10.9, 13.6, 9.7, 9.3]

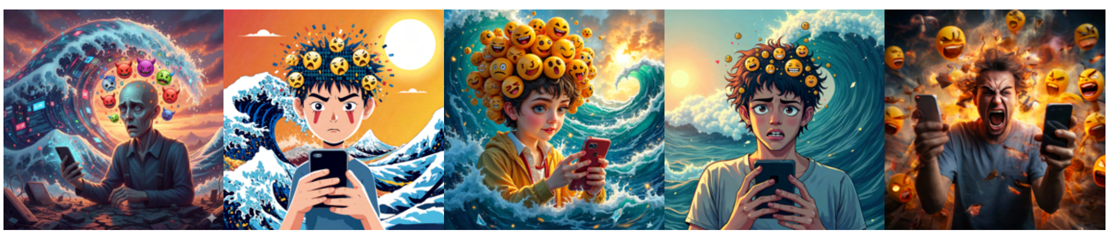

Internet Overreliance Someone holding a phone using their weak fingers, with their head surrounded by internet evil emojis, and their face showing a drained expression. The background is chaos Tsunami made from data, almost engulfing the subject. The sunlight has a clashing color. [12.7, 11.0, 11.3, 9.8, 9.5]

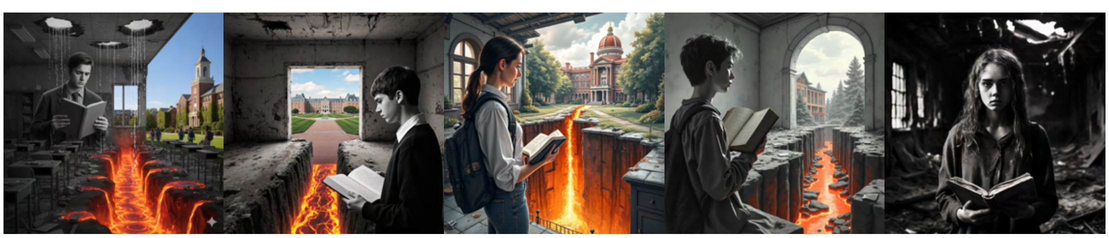

Education Inequality A college student holding a book in an old and broken classroom with a leaky roof, shows in black and white on the student's side. However, the book has no pages but only the frame. In front of them, there is a wide and deep trench with lava in it. The student shows a sad face looking at a fancy college campus on the other side of the trench, which is colorful. [9.6, 14.5, 13.5, 11.1, 12.1]

Figure 11. Social critics image images (Nano Banana, Flux Krea, Dance Flux, Flux Dev, and Playground) and corresponding HPSv3 scores.

Figure 12. (Left) Comparison between VSF-enabled Flux family models and Nano Banana. The images from the first row are from Nano Banana, and the images from the second row are from VSF applied onto Flux Krea. The two scores shown on top of the image are the HPSv3 scores rated with the wide-spectrum and original prompts. (Right) Sad emotion images generated by Flux Krea and DanceFlux. DaceFlux reinforced toxic positivity by generating a happy face.

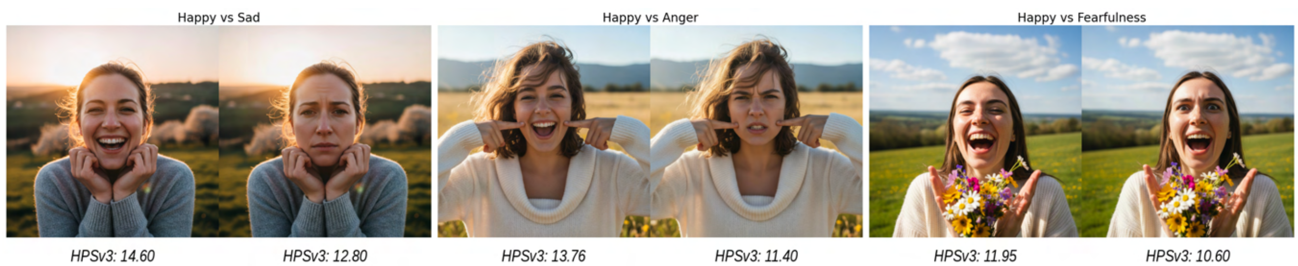

Figure 13. Emotion Bias Rating by HPSv3: all images were rated using prompts describing negative emotions, yet HPSv3 consistently assigned higher scores to the positive emotion images.

Figure 14. LoRa scale from small (left) to large (right), it gives user complete control on how “anti-aesthetics” they want their image to be. Prompt is: Night city street in the rain with pedestrians and neon signs, but the color palette is unpleasant and clashing. Everything is melting and dripping. Everything is like in a dream.

Figure 15. Top rows of images are generated with a conventional high-quality prompt, while the bottom row is generated with a wide-spectrum prompt. The left two columns are from GPT-5 Image, and the right two columns are from Nano Banana. It shows that these SOTA models can generate both conventially beauitiful images while also wide-spectrum aesthetics images.